

Noah Dimitriadis

St. Catharines, Ontario | nd22zr@brocku.ca | (289) 763-5224

Computer science student and platform engineer who builds and operates the deployment infrastructure AI products run on — GitOps release pipelines, Kubernetes and OpenShift clusters, Terraform-managed cloud resources, and LLM observability tooling — with a focus on making production systems safe to change and faster for other engineers to ship on.

Education

Brock University — Honours B.Sc. Computer Science

St. Catharines, ON | Class of 2026

- Relevant Coursework: Informed Searching Techniques, Genetic Algorithms, Neural Networks, Relational Database Systems

Professional Experience

IBM — Software Developer Co-op

Markham, ON | Feb 2026 – Present

- Created and maintain the GitOps deployment platform for the company's flagship AI product (58% of commits to the deployment/infra repo, 93 PRs, ~27K lines): 8 Helm charts and ~1,500 lines of CI/CD across three OpenShift clusters, running change detection, Terraform plan/apply, per-environment diffs, and approval-gated production deploys
- Designed and built PR preview environments — each pull request provisions an isolated namespace, three seeded Postgres databases, scoped Elasticsearch indices, and a TLS-terminated URL, with automated teardown on close; engineered the cross-instance dispatch (GitHub.com → GitHub Enterprise) so external CI failures surface on the originating PR
- Executed a live production MongoDB migration under active traffic via Terraform and the IBM Cloud CLI, coordinating credential regeneration and staging the cutover through lower environments first, with zero data loss
- Led the Kubernetes-to-OpenShift migration for the Context Manager and Core Services applications — routes replacing ingress, SecurityContextConstraints, non-root containers, scoped service accounts, and pgbouncer connection pooling in front of managed Postgres
- Built the LLM observability and evaluation platform: authored MLflow and OpenTelemetry Collector Helm charts with SSO sidecars and API-key gating, instrumented GenAI spans in the Rails application, and shipped a weekly automated eval pipeline publishing results to MLflow
- Own database provisioning and network exposure across a Terraform/Terragrunt estate spanning three environments; hardened clusters with default-deny network policies, private-only endpoints behind VPN, and SOPS-encrypted secrets, and patched a SQL injection vulnerability in a production API endpoint

Greenstone Financial Services — Junior Developer

North York, ON | Jan 2024 – Aug 2024

- Built Python automation for agent phone management, cutting manual task time by 91% while reducing errors
- Secured automation scripts in a company-wide toolkit using AWS Secrets Manager and Okta authentication
- Directed collaboration with an outside vendor to deliver new data gamification and visualization features

Personal Projects

Self-Hosted Home Server — Ubuntu, Docker, Caddy, Cloudflare Tunnel, Tailscale | Sept 2025 – Present

- Run a single-host Ubuntu server (repurposed desktop, GTX 1060) hosting a Dockerized stack behind Caddy as reverse proxy: a GPU-accelerated RAG/MCP assistant with a ChromaDB vector store, a portfolio site, a fantasy-football data-ingest pipeline, and a shared PostgreSQL instance
- Eliminated inbound port-forwarding using Cloudflare Tunnel — only a portfolio site and one Bearer-token-gated MCP service are publicly reachable; SSH has no internet-facing port, reachable only over LAN or a Tailscale tailnet
- Hardened SSH to key-only authentication with no password fallback; fail2ban rate-limits failed logins (3 attempts/10 min → 1-hour ban, escalating to a week-long ban for repeat offenders), and nightly PostgreSQL backups run automated
- Instrumented uptime monitoring as a dead-man's switch — a cron-scheduled curl heartbeat to healthchecks.io every 5 minutes, alerting in parallel via ntfy push and email on a missed check-in

WorkshopFFL Draft & Season Assistant — FastAPI, PostgreSQL, Polars | Aug 2026 – Present

- Built a fantasy football data platform (~12K lines) ingesting ten seasons of nflverse data alongside live Sleeper league state, using APScheduler for nightly ingest and 2-3 second polling during a live draft
- Designed the PostgreSQL schema and Alembic migrations, and a valuation engine that recommends picks from current roster construction and the probability a player survives to the next pick

Technical Skills

Languages: Python, Ruby, TypeScript, C#, Java, SQL

Cloud & Infrastructure: Kubernetes, OpenShift, Helm/Helmfile, Terraform/Terragrunt, Docker, GitHub Actions, IBM Cloud, SOPS, Linux, Cloudflare Tunnel, Tailscale

Backend & Data: Rails, FastAPI, React, PostgreSQL, MongoDB, Elasticsearch, Redis, Celery, SQLite

AI & Observability Tooling: OpenTelemetry, Instana, Prometheus, Grafana, Langfuse, MLflow, promptfoo, RAG (ChromaDB), Ollama, k6

Other: Git, Azure DevOps